



先进计算与数字工程研究中心

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

PATTERN RECOGNITION AND MACHINE LEARNING



Lecture 03 – Generative Classifiers and Bayesian Decisions

Chizhi Chris ZHANG

zhangchizhi@ciomp.ac.cn

Advanced Computing and Digital Technology Research Center

University of Chinese Academy of Sciences

Fall 2026

What we will build in this lecture

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Question	Tool we need
What does an alarm tell us?	Bayesian decision criterion
Which mistake matters more?	Minimum-risk decision
How can several clues be combined?	Naive Bayes classifier
Where do probabilities come from?	Parameter estimation
How do continuous features form a class?	Gaussian classification
How do we fit the class shapes?	Mean and covariance estimation



Finish line. Fit a classifier, explain its boundary, and choose an action.

Separate the state, the observation, and the action

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Object	In the factory	Notation
State	Defective or sound	Y
Observation	Alarm or no alarm	X
Action	Release, quarantine, or review	a



One important distinction

An alarm is observed. A defect is still unknown. Quarantine is a choice.

Count the four possible outcomes

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Start with the population

200 defective parts; 9,800 sound parts.

Apply the sensor rates

$200 \times 0.90 = 180$ detected defects.

$9,800 \times 0.05 = 490$ false alarms.

The other two cells. 20 missed defects; 9,310 sound parts without an alarm.

Inspection of 10,000 parts (teaching counts)

180

Defective
+ alarm

490

Sound
+ alarm

20

Defective
+ no alarm

9,310

Sound
+ no alarm

An alarm has two possible explanations


ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Joint probability

$$P(D, A) = 0.02 \times 0.90 = 0.018$$

$$P(S, A) = 0.98 \times 0.05 = 0.049$$

Read the paths. Multiply along a path. Add the paths that end at the same observation.




Only 180 of the 670 alarms are defects

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

$$P(D | A) = \frac{180}{180 + 490} = \frac{18}{67} \approx 0.269$$

The denominator has changed

We are now looking only at the alarm group, not at all defective parts.



Bayes' rule keeps the competing explanations

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

From likelihood to posterior

$$P(Y = k | x) = \frac{p(x | Y = k)P(Y = k)}{\sum_j p(x | Y = j)P(Y = j)}$$

Term	Question it answers
Prior $P(Y = k)$	How common is this class before this observation?
Likelihood $p(x Y = k)$	How compatible is the observation with this class?
Posterior $P(Y = k x)$	How plausible is this class after the observation?



Read the condition before reading the number

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Statement	Denominator	Value
$P(A D)$	All defective parts	$180/200 = 0.90$
$P(D A)$	All alarmed parts	$180/670 \approx 0.269$
$P(A)$	All inspected parts	$670/10,000 = 0.067$



Not interchangeable

$P(A | D) \neq P(D | A)$ in general.

Why the largest posterior minimizes errors

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Equal penalties for all mistakes

$$R(k | x) = \sum_{j \neq k} P(Y = j | x) = 1 - P(Y = k | x)$$

For an alarmed part

Predict defective: error probability 0.731.

Predict sound: error probability 0.269.

A restricted objective

This minimizes the number of mistakes, not every possible consequence.

Compute the posterior with three array operations

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

```
import numpy as np
prior = np.array([0.02, 0.98]) # defective, sound
alarm_rate = np.array([0.90, 0.05])
weights = prior * alarm_rate
posterior = weights / weights.sum()
print(weights)                # [0.018 0.049]
print(np.round(posterior, 3))  # [0.269 0.731]
assert np.isclose(posterior.sum(), 1.0)
```

Keep the order fixed. Every vector uses [defective, sound].

No alarm does not mean no defect

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

The remaining group

20 defective parts and 9,310 sound parts have no alarm.

$$P(D | \neg A) = \frac{20}{20 + 9,310} \approx 0.00214$$



Small is not zero

About 0.214% of the no-alarm group is defective in this example.

A different factory can produce a different answer

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Keep the sensor unchanged

$$P(A | D) = 0.90, P(A | S) = 0.05.$$


Change the prevalence

2% defective: posterior $\approx 26.9\%$.

20% defective: posterior $\approx 81.8\%$.

Check the deployment population

A training class balance is not automatically the future class prior.



Odds separate evidence from the prior

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Posterior odds

$$\frac{P(D | A)}{P(S | A)} = \frac{P(A | D)}{P(A | S)} \frac{P(D)}{P(S)}$$

Our alarm

$$18 \times \frac{1}{49} = \frac{18}{49}$$

Meaning. The alarm favors a defect by a factor of 18, but the prior odds favor sound parts by 49 to 1.

More than two classes: normalize all competitors

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Class	Prior	Likelihood	Product
A	0.40	0.30	0.12
B	0.30	0.30	0.09
C	0.30	0.1333...	0.04

$$P(Y | x) = (0.48, 0.36, 0.16)$$



Two different tasks. Choose A using the largest product. Divide by 0.25 if probabilities are requested.

Even the best rule can make mistakes

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

At the same observed feature value

Suppose $P(Y = A | x) = 0.6$ and $P(Y = B | x) = 0.4$.

Best under zero-one loss

Always choosing A gives conditional error 0.4.

What "best" means

Best under the stated distributions, available features, and loss. Not perfect; not guaranteed for a fitted model.

Try one calculation without the factory numbers

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

A new teaching example

$P(C) = 0.10$, $P(+ | C) = 0.80$, $P(+ | \neg C) = 0.10$.

Find $P(C | +)$.

Check your denominator

$$\frac{0.10 \times 0.80}{0.10 \times 0.80 + 0.90 \times 0.10} = \frac{8}{17} \approx 0.471$$

Explain the result

80% is a likelihood, not the requested posterior.

Would you release an alarmed part?

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Return to the factory

After an alarm, $P(D \mid A) = 18/67 \approx 0.269$.

The consequences

Releasing a defective part costs 30 units.

Quarantining a sound part costs 2 units.

Correct actions cost 0 in this simplified model.

A new question

Which action has the smaller expected loss?

Read a loss matrix by action and true state

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Chosen action	True defective	True sound
Quarantine	0	2
Release	30	0



Convention used throughout this lecture

Rows are actions. Columns are true states.

$L(a, j)$ is the loss if we choose action a and state j is true.

Read one cell aloud. Release + true defective = a missed defect, loss 30.

Conditional risk is a weighted average of losses

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

$$R(a | x) = \sum_j L(a, j) P(Y = j | x)$$

Quarantine

$$R(Q | A) = 0 \times P(D | A) + 2 \times P(S | A).$$

Release

$$R(R | A) = 30 \times P(D | A) + 0 \times P(S | A).$$

Decision rule. Choose $a^*(x) = \arg \min_a R(a | x)$.

Lower probability can still justify quarantine

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Action	Exact expected loss	Approximation
Quarantine	$2 \times 49/67$	1.463
Release	$30 \times 18/67$	8.060



Decision

Quarantine has the lower expected loss.

Do not rewrite the probability

The posterior remains 0.269. Only the action rule changed.

Derive the decision threshold from the losses

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Let $p = P(D | x)$

$$\underbrace{2(1-p)}_{\text{quarantine}} < \underbrace{30p}_{\text{release}}$$

$$p > \frac{2}{30+2} = 0.0625$$




General zero-correct-loss case

$$p > \frac{c_{\text{false alarm}}}{c_{\text{miss}} + c_{\text{false alarm}}}$$

The threshold is where the risk curves cross

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER



Read the graph. Release costs $30p$; quarantine costs $2(1 - p)$. The lower curve determines the action.

At an alarm. $p \approx 0.269$ is to the right of the crossing: quarantine.

Changing a cost changes the policy, not the sensor

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Miss cost	False-alarm cost	Quarantine threshold
30	2	0.0625
30	10	0.25
90	2	0.0217



Interpretation

Expensive false alarms require stronger evidence. Expensive misses require earlier intervention.

A third action: ask for a review

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Action	Defective	Sound
Quarantine	0	2
Release	30	0
Review	0.5	0.5



For the alarmed part


Risks: quarantine 1.463, release 8.060, review 0.5.
Choose review under this stated loss model.

A modelling assumption

The review cost includes all relevant consequences here; real review can have delays and errors.

Review occupies the uncertain middle region

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER



Three actions. Release for $p < 1/60$; review for $1/60 < p < 0.75$; quarantine for $p > 0.75$.

At the boundaries. Two actions tie. Pick a consistent tie policy.

The most likely class need not give the best action

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Action	True A	True B	True C	Risk
Choose A	0	4	8	2.20
Choose B	2	0	3	1.65
Choose C	5	1	0	3.25



Posterior probabilities

$$P(Y | x) = (0.60, 0.25, 0.15).$$

Class A is most likely; choosing B has the lowest expected loss.

Turn a loss table into a matrix product

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

```
post_three = np.array([0.60, 0.25, 0.15])
loss_three = np.array([[0., 4., 8.],
                        [2., 0., 3.],
                        [5., 1., 0.]])
risks_three = loss_three @ post_three
print(risks_three)           # [2.2  1.65 3.25]
print(risks_three.argmin()) # 1: the second action
```

Check the meaning of the index

Python index 1 means the second row, not class label 1 in every dataset.

An everyday example: keep or drop a course?

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

A separate teaching example

$P(\text{good}) = 0.6$. An engaging class occurs with probability 0.8 in good courses and 0.1 in poor courses.

After one engaging class

$$P(\text{good} \mid \text{engaging}) = \frac{0.48}{0.52} = \frac{12}{13}$$

Losses. Keep a poor course: 10. Drop a good course: 20. Correct choices: 0. Risks are 10/13 versus 240/13; keep it.

Pointwise decisions minimize overall expected risk

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Average over future observations

$$\mathcal{R}(a) = E_X[R(a(X) | X)], \quad \mathcal{R}^* = E_X[\min_a R(a | X)]$$

Factory model, with only release or quarantine

Quarantine alarms and release non-alarms. Expected loss per part:

$$\frac{490 \times 2 + 20 \times 30}{10,000} = 0.158$$

Bayes risk. The minimum under the true distribution, specified loss, and available actions. An estimated model need not attain it.

Can a few words reveal the kind of message?

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Two messages

One promises a free prize. One arranges a project meeting.

The task

Predict spam or ordinary mail from the words that appear.

A practical question. How can we assign a probability to a whole message?

Two messages, one classification problem

Message A

FREE prize! Claim your reward.

Message B

Project meeting at 10. Bring your draft.

Turn words into explicit observations

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Message fragment	free present?	meeting present?	free count
free prize	1	0	1
free free reward	1	0	2
project meeting	0	1	0
free meeting room	1	1	1



Choose a representation

Presence and counts answer different questions. Do not switch between them silently.

The simplifying assumption sits inside each class

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER


Naive Bayes assumption

$$p(x_1, \dots, x_d \mid Y = k) = \prod_{j=1}^d p(x_j \mid Y = k)$$

Read the graph. Condition on the class at the top.
Model each observed feature separately.

Not a statement about language

Real words need not be independent; this is the model assumption.



Why the simplification saves so many parameters

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Ten binary features, within one class	Free probabilities
Full joint table: 2^{10} possible patterns	$2^{10} - 1 = 1,023$
Independent Bernoulli features	10



With two classes

2,046 class-conditional parameters become 20, plus one class-prior parameter.

The trade-off

Fewer parameters mean less to estimate, but less freedom to represent interactions.

Conditional independence is not marginal independence

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

A small probability model

Equal class priors. In class A, both features have probability 0.9; in class B, both have probability 0.1. They are independent within each class.

$$P(X_1 = 1) = P(X_2 = 1) = 0.5$$
$$P(X_1 = 1, X_2 = 1) = 0.5(0.9^2) + 0.5(0.1^2) = 0.41$$



After mixing classes

$0.41 \neq 0.5 \times 0.5$: the features are not independent overall.

Learn presence probabilities from labelled messages

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Training counts	Spam	Ordinary
Number of messages	4	6
Messages containing free	3	1
Messages containing meeting	0	4



Before smoothing

$$P(\text{spam}) = 4/10.$$

$$P(\text{free} = 1 \mid \text{spam}) = 3/4.$$

Observation model. One Bernoulli variable per word: present or absent.

An unseen event should not erase every other clue

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

The zero-frequency problem

No spam message in four training examples contains meeting. A raw estimate gives probability 0.

Add-one smoothing for a binary feature

$$\hat{\phi}_{kj} = \frac{n_{kj} + 1}{n_k + 2}$$

For meeting in spam

$(0 + 1)/(4 + 2) = 1/6$, not 0.

An absent word is evidence in the Bernoulli model

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Smoothed probability	Spam	Ordinary
free present	4/6	2/8
meeting present	1/6	5/8



Message: free present, meeting absent

$$w_S = 0.4 \times \frac{4}{6} \times \left(1 - \frac{1}{6}\right) \approx 0.2222$$

$$w_O = 0.6 \times \frac{2}{8} \times \left(1 - \frac{5}{8}\right) = 0.05625$$

A complete Bernoulli prediction in code

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

```
class_count = np.array([4., 6.])
present = np.array([[3., 0.], [1., 4.]])
phi_binary = (present + 1) / (class_count[:, None] + 2)
x_binary = np.array([1., 0.]) # free, meeting
likelihood_binary = (phi_binary ** x_binary *
                    (1 - phi_binary) ** (1 - x_binary)).prod(axis=1)
w_binary = class_count / class_count.sum() * likelihood_binary
print(np.round(w_binary / w_binary.sum(), 3)) # [0.798 0.202]
```

Two dimensions to track. Rows are classes; columns are words. The input describes one new message.

Counts use a different sample space

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Training word counts	free	meeting	project	Total
Spam	6	0	2	8
Ordinary	1	5	4	10



Add-one smoothing for a vocabulary of size 3

$$\hat{\phi}_{kj} = \frac{n_{kj} + 1}{\sum_v n_{kv} + 3}$$

New example. These are token counts, not the message-presence counts on the previous slides. Class priors remain 0.4 and 0.6.

Score the message "free free project"

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Feature vector

$x = (2, 0, 1)$ in vocabulary order [free, meeting, project].

$$w_S = 0.4 \left(\frac{7}{11} \right)^2 \left(\frac{3}{11} \right) \approx 0.04418$$


$$w_O = 0.6 \left(\frac{2}{13} \right)^2 \left(\frac{5}{13} \right) \approx 0.005462$$



A common factor cancels. The multinomial coefficient is identical for both classes and is omitted when comparing scores.

Use log scores when the word list gets long

ADVANCED CO



$$s_k = \log P(Y = k) + \sum_j x_j \log \phi_{kj}$$

Read the contributions. Repeated free contributes twice; absent meeting contributes zero. Add log probabilities to avoid tiny products.

Normalize log scores without numerical overflow

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

```
token_counts = np.array([[6., 0., 2.], [1., 5., 4.]])
phi_count = (token_counts + 1) / (token_counts.sum(axis=1, keepdims=True) + 3)
x_count = np.array([2., 0., 1.])
log_scores = np.log([0.4, 0.6]) + np.log(phi_count) @ x_count
shifted = log_scores - log_scores.max()
post_count = np.exp(shifted) / np.exp(shifted).sum()
print(np.round(post_count, 3)) # [0.89 0.11]
```

Subtracting the same constant is safe

It preserves all score differences and cancels during normalization.

Copying a feature can double-count the evidence

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

One clue

Prior odds 1 : 1 and likelihood ratio 2 : 1 give posterior odds 2 : 1, or probability 2/3.

Duplicate the same clue as a new feature

A naive product gives odds 4 : 1, or probability 4/5. No new information was observed.

The problem

Two copies are perfectly dependent within a class. The independence assumption is violated.

A bag of words does not preserve who did what

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Sentence A

The dog bites the man.

Sentence B

The man bites the dog.

count vector A = count vector B

An information limit

Any classifier using only this identical count vector must treat the two sentences identically.

Not every discrete feature is a word count

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Account feature	Possible values	Distribution
Posting frequency	low, medium, high	Categorical
Profile photo supplied	yes, no	Bernoulli
Number of posts by topic	counts across topics	Multinomial



Categorical smoothing

$$P(X_j = m \mid Y = k) = \frac{n_{kjm} + 1}{n_k + M_j}$$

The observation determines the distribution family

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

What is measured?	Possible model	Parameters within each class
Present or absent	Bernoulli	One probability per feature
One of several categories	Categorical	A probability vector
Counts of vocabulary tokens	Multinomial	Word probabilities
Continuous measurement	Gaussian	Mean vector and covariance matrix



Type is only a starting point

A continuous feature is not automatically Gaussian. Check whether the assumption is plausible.

What does training actually have to estimate?

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

For the models we just built

- Class priors from labelled messages.
- Feature probabilities from within-class counts.
- A vocabulary fixed using training data.



A question for the next section

If four messages are all spam, should the estimated spam probability be exactly 1?

Keep evaluation separate

Test messages must not be used to build counts or choose smoothing settings.

A model family is not yet a fitted model

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Object	Example	What is unknown?
Observation X	A binary outcome, 0 or 1	The next outcome
Parameter θ	$P(X = 1) = \theta$	The probability generating outcomes
Data $\mathcal{D}$	Observed past outcomes	Nothing: these are now fixed records

Two fitting ideas

Find the parameter that best explains the data, or update a probability distribution over the parameter.

Ten outcomes are evidence about one probability

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

A new estimation example

 $\mathcal{D} = (1, 1, 0, 1, 0, 1, 1, 1, 0, 1).$

Seven successes and three failures, independently sampled with a common θ .

Likelihood of this observed sequence

$$L(\theta) = P(\mathcal{D} \mid \theta) = \theta^7 (1 - \theta)^3$$

Keep the interpretation straight. The data are fixed; the parameter is varied.

The likelihood peaks near the observed fraction

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Compare candidate parameters

Too small a θ makes seven successes unlikely.


Too large a θ makes three failures unlikely.

The peak

$$\hat{\theta} = 7/10 = 0.7.$$

Vertical axis

Likelihood of the observed sequence, not a density over the parameter.



Maximum likelihood gives the sample proportion

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

$$\ell(\theta) = 7 \log \theta + 3 \log(1 - \theta)$$

$$\frac{d\ell}{d\theta} = \frac{7}{\theta} - \frac{3}{1 - \theta} = 0$$

$$7(1 - \theta) = 3\theta \quad \Rightarrow \quad \hat{\theta} = 0.7$$



General result

For s successes in n independent Bernoulli trials, $\hat{\theta} = s/n$.

Three successes do not prove certainty

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Only three observations

$\mathcal{D} = (1, 1, 1)$ gives $L(\theta) = \theta^3$ and $\hat{\theta} = 1$.

What the estimate says

1 is the maximizer of the likelihood, not proof that failures are impossible.

Why we need another view

A point estimate hides how much evidence supports it.

Class posterior and parameter posterior differ

Earlier classification

$$P(Y = k | x)$$

Which class produced this observation?

A distribution across class labels


Now parameter estimation

$$p(\theta | \mathcal{D})$$

Which parameter values explain the data?

A distribution across possible parameters

$$p(\theta | \mathcal{D}) \propto p(\mathcal{D} | \theta) p(\theta)$$



A prior describes uncertainty before this dataset

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER


A Beta prior for a probability

$$p(\theta) \propto \theta^{\alpha-1}(1-\theta)^{\beta-1}, \quad 0 < \theta < 1$$

Three examples. Beta(1,1): uniform. Beta(2,2): centered. Beta(8,2): favors larger probabilities.


A choice to justify

A convenient prior is not automatically a suitable prior.



Updating Beta counts gives a full posterior

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER



$$p(\theta | \mathcal{D}) \propto \theta^8 (1 - \theta)^4 \Rightarrow \theta | \mathcal{D} \sim \text{Beta}(9, 5)$$

General update. Beta(α, β) becomes Beta($\alpha + s, \beta + n - s$).

Three estimates answer three different questions

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Estimate	Formula in this example	Value
Maximum likelihood	$7/10$	0.700
Posterior mean	$9/(9 + 5)$	0.643
Posterior mode	$(9 - 1)/(9 + 5 - 2)$	0.667



Mode formula condition

$(\alpha - 1)/(\alpha + \beta - 2)$ applies to an interior Beta mode when both parameters exceed 1.

Prediction averages over parameter uncertainty

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

$$P(X_{\text{new}} = 1 \mid \mathcal{D}) = \int_0^1 \theta p(\theta \mid \mathcal{D}) d\theta = \frac{9}{14}$$

One future Bernoulli outcome

Its predictive success probability equals the posterior mean of the parameter.

Do not overgeneralize

For nonlinear predictions, averaging predictions need not equal prediction at the average parameter.

More data reduce the influence of a fixed prior

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER


$$E[\theta | \mathcal{D}] = \frac{n}{n+4} \frac{s}{n} + \frac{4}{n+4} \frac{1}{2}$$

Keep the observed success fraction at 0.7

$n = 10$: $9/14 \approx 0.643$.

$n = 100$: $72/104 \approx 0.692$.

Meaning. With a fixed Beta(2,2) prior, the data receive more weight as the sample size grows.



Updating in batches should give the same answer

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Stage	Successes added	Failures added	Posterior
Start	0	0	Beta(2,2)
First batch	3	1	Beta(5,3)
Second batch	4	2	Beta(9,5)



Compare with one batch

All seven successes and three failures at once also give Beta(9,5).

Do not count the first batch twice

The first posterior already contains the first batch.

Check a Bayesian update with a short program

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

```
alpha_prior, beta_prior = 2., 2.
successes, failures = 7, 3
alpha_post = alpha_prior + successes
beta_post = beta_prior + failures
post_mean = alpha_post / (alpha_post + beta_post)
post_mode = (alpha_post - 1) / (alpha_post + beta_post - 2)
print(round(post_mean, 6), round(post_mode, 6))
# 0.642857 0.666667
assert (2 + 3 + 4, 2 + 1 + 2) == (alpha_post, beta_post)
```

No fitting loop is needed here. The conjugate update has a closed form.

The same counting idea extends to several outcomes

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Three categories

Observed counts (3, 1, 0). Maximum-likelihood probabilities: (3/4, 1/4, 0).

Uniform Dirichlet prior

$$\theta \sim \text{Dirichlet}(1, 1, 1)$$

$$\theta \mid \mathcal{D} \sim \text{Dirichlet}(4, 2, 1)$$

Posterior mean. (4/7, 2/7, 1/7). Add one to each count, then normalize.

A Gaussian mean can also be learned with a prior

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Known observation variance

$X_i \mid \mu \sim \mathcal{N}(\mu, \sigma^2)$, independently; $\mu \sim \mathcal{N}(m_0, v_0)$.

$$v_n = \left(v_0^{-1} + n/\sigma^2 \right)^{-1}, \quad m_n = v_n \left(m_0/v_0 + n\bar{x}/\sigma^2 \right)$$



A numerical update

$m_0 = 0, v_0 = 4, \sigma^2 = 4, n = 4, \bar{x} = 3$.

Posterior for μ : $\mathcal{N}(2.4, 0.8)$.

A sharper parameter posterior is not noiseless data

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

What the original illustration shows


Posterior curves for an unknown mean become more concentrated as observations accumulate.

Two different variances

Uncertainty about the mean can shrink while individual measurements still vary.

Prediction with known noise

variance. $\text{Var}(X_{\text{new}} | \mathcal{D}) = \sigma^2 + v_n$, not just v_n .



One last check: all successes with a uniform prior

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Three successes, no failures

$\theta \sim \text{Beta}(1, 1)$ before the observations.

Quantity	Answer
Posterior	Beta(4,1)
Maximum-likelihood estimate	1
Posterior mean and next-success probability	4/5
Posterior mode	1, at the boundary



Next. For continuous feature vectors, we need a model for both the centre and the shape.

Continuous measurements form clouds, not word lists

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER


Two measurements per part

$\mathbf{x} = (x_1, x_2)^T$: vibration and temperature in scaled units.

A generative model for each class

$$p(\mathbf{x} \mid Y = k) = \mathcal{N}(\mathbf{x}; \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k).$$

Question. Which class distribution gives stronger support to a new point?



A two-dimensional density is a surface

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Read the three axes


x_1 and x_2 describe the feature location. Height is the density $p(\mathbf{x})$.

A probability is a volume

Integrate the density over a region of the feature plane.

Height is not class probability

A high density means local concentration within this distribution, not certainty about a class.



A density can exceed one without any contradiction

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

$$p(x | \mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp \left[-\frac{(x - \mu)^2}{2\sigma^2} \right]$$



A narrow distribution

For $X \sim \mathcal{N}(0, 0.01)$, the standard deviation is 0.1 and $p(0) \approx 3.989$.

Probability stays between zero and one

$P(-0.1 \leq X \leq 0.1) \approx 0.683$. It is an area, not the height 3.989.

The same centre can hide different class shapes



$$\mu = \begin{bmatrix} 3 \\ 2 \end{bmatrix}, \quad \Sigma_{\text{left}} = \begin{bmatrix} 25 & 0 \\ 0 & 9 \end{bmatrix}, \quad \Sigma_{\text{right}} = \begin{bmatrix} 15 & 5 \\ 5 & 5 \end{bmatrix}$$

Compare geometry. The left density is axis-aligned. The positive off-diagonal covariance on the right tilts the contours upward.

Distance should respect the shape of the cloud

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER


Squared Mahalanobis distance

$$d_M^2 = (\mathbf{x} - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu})$$

With $\boldsymbol{\mu} = 0$ and $\boldsymbol{\Sigma} = \text{diag}(4, 1)$

$A=(2, 0)$ and $B=(0, 1)$ both have $d_M^2 = 1$. Their Euclidean distances are 2 and 1.

Equal Mahalanobis distance, unequal Euclidean distance



Three ingredients determine Gaussian density

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

$$p(\mathbf{x} | k) = \frac{\exp[-\frac{1}{2}(\mathbf{x} - \mu_k)^T \Sigma_k^{-1}(\mathbf{x} - \mu_k)]}{(2\pi)^{d/2} |\Sigma_k|^{1/2}}$$

Ingredient	Role
Mean μ_k	Centre of the class
Inverse covariance Σ_k^{-1}	Shape-aware distance
Determinant $ \Sigma_k $	Normalizes the spread and height



Assumption for this density formula

The covariance matrix is symmetric and positive definite.

Use a log score to compare Gaussian classes

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Class score, dropping a common constant

$$g_k(\mathbf{x}) = -\frac{1}{2} \log |\Sigma_k| - \frac{1}{2} (\mathbf{x} - \mu_k)^T \Sigma_k^{-1} (\mathbf{x} - \mu_k) + \log \pi_k$$

Three terms. Spread penalty + shape-aware distance penalty + log prior.



What may be removed

$-d \log(2\pi)/2$ is shared by all classes. A class-specific determinant is not.


Different covariance matrices can bend the boundary

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

The plotted example. $\mu_A = (0, 0)^T$, $\mu_B = (3, 2)^T$;
equal priors. Covariances differ.

Boundary equation. $g_A(\mathbf{x}) = g_B(\mathbf{x})$. The quadratic
terms generally do not cancel.

Class density contours and decision boundary



Shared covariance makes the boundary linear

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

$$\Sigma_A = \Sigma_B = \Sigma$$

$$g_k(\mathbf{x}) = \mathbf{x}^T \Sigma^{-1} \mu_k - \frac{1}{2} \mu_k^T \Sigma^{-1} \mu_k + \log \pi_k + \text{common terms}$$



After comparing two classes

$$g_A(\mathbf{x}) - g_B(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b$$

Why. The same quadratic term and the same log determinant cancel.


Read contours and the separating line together

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

In this model. Both classes share covariance with positive correlation and have equal priors.

What remains different. Their centres. The black boundary is straight even though each class has elliptical contours.

Class density contours and decision boundary



Nearest mean is a special case, not the whole rule


ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

A sufficient set of conditions

Zero-one loss, equal priors, and one shared spherical covariance: $\Sigma_k = \sigma^2 I$.

$$\arg \max_k g_k(\mathbf{x}) = \arg \min_k \|\mathbf{x} - \mu_k\|^2$$

Change the assumption	What changes?
Shared but non-spherical covariance	Use a shape-aware metric
Unequal priors	Add a class preference
Different class covariances	Distance and volume both matter



Priors move a boundary even when densities stay fixed

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

One-dimensional example


$X | A \sim \mathcal{N}(0, 1); X | B \sim \mathcal{N}(4, 1).$

$$x^* = 2 + \frac{1}{4} \log \frac{\pi_A}{\pi_B}$$

Change the priors

Equal priors: $x^* = 2.$

$(\pi_A, \pi_B) = (0.8, 0.2): x^* \approx 2.347.$



One class can win in both tails


Same mean, unequal variances

$X | A \sim \mathcal{N}(0, 1)$, $X | B \sim \mathcal{N}(0, 4)$; equal priors.

$$|x| = \sqrt{\frac{8}{3} \log 2} \approx 1.360$$

Decision regions

Choose A near 0; choose B in both tails.



Diagonal covariance does not guarantee a line

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Gaussian Naive Bayes

Each class has a diagonal covariance matrix; features are independent within that Gaussian class.

Class-specific variances

If the diagonal entries differ between classes, quadratic terms can remain.

Different restrictions

Diagonal within each class $\neq$ shared across classes. Shared covariance is the condition used for a linear boundary.

Inspect every term in a numerical Gaussian score

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Given parameters

$\mu_A = (1, 1)^T$, $\mu_B = (5, 5)^T$, $\Sigma_A = \Sigma_B = I$, equal priors. New point: $(2, 3)^T$.

Term	Class A	Class B
Squared distance	5	13
Log determinant	0	0
Score without common normalizer	$-2.5 + \log 0.5$	$-6.5 + \log 0.5$



$$P(A | \mathbf{x}) = \frac{1}{1 + e^{-4}} \approx 0.982$$

Compute a Gaussian score without forming an inverse

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

```
def gaussian_score(x, mean, covariance, prior):  
    delta = x - mean  
    sign, logdet = np.linalg.slogdet(covariance)  
    np.linalg.cholesky(covariance) # requires positive definiteness  
    if sign <= 0 or prior <= 0:  
        raise ValueError("Invalid covariance or prior")  
    distance2 = delta @ np.linalg.solve(covariance, delta)  
    return -0.5 * (logdet + distance2) + np.log(prior)
```

Input contract. A symmetric positive-definite covariance and a positive prior. The function returns an unnormalized log score.

Start from eight points, not from given parameters

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Class A


$(0, 0), (2, 0), (2, 2), (0, 2)$.

Class B

$(4, 4), (6, 4), (6, 6), (4, 6)$.

Our complete task

Estimate the parameters, derive the boundary, then classify new points.



Estimate one mean and one prior per class

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

$$\hat{\mu}_k = \frac{1}{n_k} \sum_{i:y_i=k} \mathbf{x}_i, \quad \hat{\pi}_k = \frac{n_k}{n}$$

The eight-point dataset

$$\hat{\mu}_A = (1, 1)^T, \hat{\mu}_B = (5, 5)^T.$$

$$n_A = n_B = 4, \text{ so } \hat{\pi}_A = \hat{\pi}_B = 0.5.$$

Keep the labels attached

Compute class means from the samples in that class, not from the entire mixed cloud.

Centre the samples before measuring their spread

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

A sample	Centred vector	Product of coordinates
(0, 0)	(-1, -1)	1
(2, 0)	(1, -1)	-1
(2, 2)	(1, 1)	1
(0, 2)	(-1, 1)	-1



One outer product

$$\begin{bmatrix} 1 \\ -1 \end{bmatrix} \begin{bmatrix} 1 & -1 \end{bmatrix} = \begin{bmatrix} 1 & -1 \\ -1 & 1 \end{bmatrix}$$

Maximum-likelihood covariance averages outer products

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

$$\hat{\Sigma}_k = \frac{1}{n_k} \sum_{i: y_i=k} (\mathbf{x}_i - \hat{\mu}_k)(\mathbf{x}_i - \hat{\mu}_k)^T$$

For class A

$$\sum_i (\mathbf{x}_i - \hat{\mu}_A)(\mathbf{x}_i - \hat{\mu}_A)^T = \begin{bmatrix} 4 & 0 \\ 0 & 4 \end{bmatrix}$$

$$\hat{\Sigma}_A = \hat{\Sigma}_B = I$$

Divide by n or n minus one? State the estimator

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Estimator	Class-A covariance	NumPy setting
Gaussian maximum likelihood	I	<code>ddof=0</code>
Unbiased sample covariance	$\frac{4}{3} I$	<code>ddof=1</code>



Same numerator, different objectives

Maximum likelihood fits the Gaussian model. The unbiased version corrects the expected covariance estimate under repeated sampling.

Defaults can differ

Use an explicit denominator when matching a derivation to code.

Shared covariance is pooled within classes

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

$$\hat{\Sigma}_{\text{shared}} = \frac{1}{n} \sum_k \sum_{i: y_i = k} (\mathbf{x}_i - \hat{\mu}_k)(\mathbf{x}_i - \hat{\mu}_k)^T$$

Correct within-class result

For these eight points, $\hat{\Sigma}_{\text{shared}} = I$.

Not the covariance around the grand mean

$$\text{Cov}_{\text{all}}(X) = \begin{bmatrix} 5 & 4 \\ 4 & 5 \end{bmatrix} \quad \text{with denominator } n.$$


Fit the parameters directly from the labelled array

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

```
X = np.array([[0.,0.],[2.,0.],[2.,2.],[0.,2.],
              [4.,4.],[6.,4.],[6.,6.],[4.,6.]])
y = np.array([0,0,0,0,1,1,1,1])
means = np.array([X[y == k].mean(axis=0) for k in (0, 1)])
covariances = np.array([np.cov(X[y == k], rowvar=False, ddof=0)
                        for k in (0, 1)])
priors = np.bincount(y) / len(y)
print(means) # [[1. 1.] [5. 5.]]
assert np.allclose(covariances, np.array([np.eye(2), np.eye(2)]))
```

Array convention. Rows are samples; columns are features. Labels are stored separately.

The fitted boundary is $x_1 + x_2 = 6$



$$\|\mathbf{x} - (1, 1)^T\|^2 = \|\mathbf{x} - (5, 5)^T\|^2$$

$$8x_1 + 8x_2 - 48 = 0 \Rightarrow x_1 + x_2 = 6$$

Regions. Choose A below the line and B above it. Points on the line tie under equal priors and zero-one loss.

Predict a new point with the fitted model

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

```
new_point = np.array([4., 3.])
new_scores = np.array([gaussian_score(new_point, means[k],
                                     covariances[k], priors[k])
                       for k in (0, 1)])
new_weights = np.exp(new_scores - new_scores.max())
new_posterior = new_weights / new_weights.sum()
print(np.round(new_posterior, 3)) # [0.018 0.982]
print(["A", "B"][new_posterior.argmax()]) # B
```

Geometry agrees. $4 + 3 > 6$, so the point lies in the class B region.

Too few independent directions make covariance singular

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Three points on one line


$(0, 0)$, $(1, 2)$, $(2, 4)$ have only one direction of variation.

$$\hat{\Sigma} = \begin{bmatrix} 2/3 & 4/3 \\ 4/3 & 8/3 \end{bmatrix}, \quad |\hat{\Sigma}| = 0$$

A general rank limit

With n_k samples in d dimensions, rank is at most $\min(d, n_k - 1)$. Even more samples can be collinear.

Centred observations confined to one direction



Shrinkage gives weak directions some variance

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER


$$\tilde{\Sigma} = (1 - \lambda)S + \lambda\tau I, \quad \tau = \text{tr}(S)/d$$

For the collinear example

$\lambda = 0.2$, $\tau = 5/3$. Eigenvalues change from $(0, 10/3)$ to $(1/3, 3)$.

A bias-stability trade-off

The added variance is a modelling choice. Choose its strength using training/validation data, not the final test set.



Covariance structure controls how much must be learned

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Structure for two classes, $d=20$	Covariance parameters
Separate full covariance matrices	$2 \times 210 = 420$
One shared full covariance matrix	210
Separate diagonal covariance matrices	$2 \times 20 = 40$
One shared spherical variance	1



These counts exclude means and priors. A symmetric full matrix has $d(d+1)/2$ free entries.

Choose with evidence

More flexible is not always better when each class has few observations.

Update means and covariance as samples arrive

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Keep a count, a mean, and a scatter matrix

$$\delta = \mathbf{x}_{n+1} - \mu_n, \quad \mu_{n+1} = \mu_n + \frac{\delta}{n+1}$$

$$M_{n+1} = M_n + \delta(\mathbf{x}_{n+1} - \mu_{n+1})^T$$



Maximum-likelihood covariance

After n samples, $\hat{\Sigma} = M_n/n$.

For the scalar sequence 2, 4, 6: mean 4, scatter 8, variance 8/3.

Per-class training. Maintain a separate state for each class; initialize count 0, mean 0, scatter 0.

Test the decision objective, not just accuracy

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

On 100 labelled test parts	Policy A	Policy B
Missed defects, out of 20	2	1
False alarms, out of 80 sound parts	8	16
Accuracy	90%	83%
Total loss: 30 per miss, 2 per false alarm	76	62



Read both outcomes

Policy A makes fewer errors. Policy B has lower total loss under these stated costs.

Teaching comparison. Invented test counts illustrate the metric trade-off; they are not results from the eight-point dataset.

Find three errors before trusting a prediction

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Broken workflow

- Estimate one covariance around the grand mean and call it within-class covariance.
- Use test labels to adjust the model until all test points are correct.
- Multiply a loss matrix by unnormalized log scores.



Repair it

- Centre within each class before pooling.
- Tune on training/validation data; reserve a final test set.
- Normalize scores into posterior probabilities before computing risks.

Finish the factory decision from data to action

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

A fresh labelled training sample

10 defective parts: 8 alarms.

90 sound parts: 18 alarms. Use unsmoothed frequency estimates.

For a new alarm

$$P(D | A) = \frac{(10/100)(8/10)}{(10/100)(8/10) + (90/100)(18/90)} = \frac{4}{13}$$

Choose an action. Miss cost 30; false-alarm cost 2; review cost 0.5. Risks: release 120/13, quarantine 18/13, review 0.5. Choose review.