

先进计算与数字工程研究中心

Chizhi Chris ZHANG

zhangchizhi@ciomp.ac.cn

Advanced Computing and Digital Technology Research Center

University of Chinese Academy of Sciences

Fall 2026

reaches the sorting gate

The observation


Color score $x_1 = 7.2$ and firmness $x_2 = 4.1$.

The decision

Pack now or hold for another day?

The complication

Two varieties overlap in the sensor space.



ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Question	Quantity	Typical use
How could this measurement arise?	$p(x y)$ and $p(y)$	Describe each class
Which label fits this measurement?	$p(y x)$	Predict a class
Which side of the boundary is it on?	$g(x)$	Make a decision



Lecture map

先进计算与数字工程研究中心

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Syllabus topic	Question answered
Classifier families	What should the model represent?
Linear discriminant functions	What does one straight boundary mean?
Generalized linear functions	How can a feature map create a curve?
Piecewise linear functions	How can several local lines form one rule?
Fisher linear discriminant	Which projection separates the classes?



Programming

NumPy examples will calculate scores, maps, prototypes, and Fisher directions.

World model and decision model

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Generative route

Fit $p(x | y)$ and $p(y)$, then derive $p(y | x)$.


Discriminative route

Fit $p(y | x)$ or a score $g(x)$ directly.

Do not rank them by name

The useful choice depends on data, assumptions, and the required output.

Generative: describe each class



$p(x | y)$ and $p(y)$

Discriminative: describe the decision



$p(y | x)$ or $g(x)$

Generative modelling describes each class

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Tomato example

Fit a color and firmness distribution for each maturity class.

$$p(y = k | x) = \frac{p(x | y = k)p(y = k)}{\sum_j p(x | y = j)p(y = j)}$$

What the model can do. Classify, sample plausible measurements, and reveal unusual observations.

Discriminative modelling targets the label

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER


Conditional probability

Estimate $p(y | x)$ when calibrated uncertainty matters.

Decision score

Estimate $g(x)$ when only the class decision is required.

$$\hat{y} = \arg \max_k p(y = k | x) \quad \text{or} \quad \hat{y} = \text{sign } g(x)$$



Bayes can produce a discriminant function

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Start from the posterior

Choose class 1 when $p(y = 1 | x) > p(y = 2 | x)$.

$$g(x) = \log p(x | y = 1) + \log p(y = 1) - \log p(x | y = 2) - \log p(y = 2)$$



Same decision

Class 1 is selected when $g(x) > 0$.

What unlabeled inputs can contribute

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Possible generative use

Unlabeled x values can help estimate the structure of $p(x)$ under a stated model.

No automatic benefit

A wrong mixture or cluster assumption can move the boundary in the wrong direction.

Purely supervised discriminative fit

Without an added objective, unlabeled rows contribute no label loss.

Fewer assumptions can focus capacity on the boundary

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

A direct target

A discriminative model need not reproduce every detail of $p(x)$.

Tomato sensor

Lighting variation that does not change the label may be ignored by the boundary.

The tradeoff

Less structure also means fewer ways to use missing labels or detect implausible inputs.

When generative assumptions help

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Situation	Why the structure helps
Few labels and credible class models	Parameters share information through the model
Missing measurements	The joint distribution can marginalize missing entries
Anomaly screening	Low likelihood can flag an unfamiliar input
Simulation is required	The fitted model can produce synthetic observations

When generative assumptions hurt

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Misspecified shape


A Gaussian ellipse cannot represent every class distribution.

Hidden groups

Two tomato varieties may create two modes inside the same maturity class.

Consequence

The estimated likelihood can be wrong far from the boundary and still alter the decision.



A fair comparison uses the same evidence

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

- 1 Use the same training, validation, and test split.
- 2 Tune both model families on the validation set.
- 3 Evaluate the metric that matches the decision cost.
- 4 Inspect calibration only when probabilities are claimed.



Common mistake

Comparing one tuned model with another model's default settings proves little.

Case change under a different light

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Morning

Color and firmness separate the two decisions with a simple diagonal boundary.

Afternoon

A warm lamp shifts the color score by $+1.5$ for both classes.

Question

Which part should change: the data model, the features, or the boundary bias?

Need	First model to consider
A direct boundary with many labels	Discriminative
A class distribution and simulated inputs	Generative
Useful unlabeled structure with credible assumptions	Generative or semi supervised
Only calibrated class probabilities	Conditional discriminative

Choose one row and state the assumption that makes your choice reasonable.

Four questions for every boundary

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

- 1 What representation enters $g(x)$?
- 2 What shape can $g(x) = 0$ express?
- 3 How does the score become a class decision?
- 4 What evidence would reveal a bad assumption?



Use this checklist

Apply the same four questions to linear, mapped, piecewise, and Fisher boundaries.

Discriminant functions


The model

A weighted sum turns a feature vector into one score.

$$g(x) = w^\top x + w_0$$

The boundary

All points satisfying $g(x) = 0$.



Use a fixed tie rule or send the sample for review.

A score is not a label

The sign creates the label. The magnitude describes confidence only under extra assumptions.

100


© 2011 Pearson Education, Inc.



For boundary points x_a and x_b , $w^\top (x_a - x_b) = 0$.

$$w \perp (x_a - x_b)$$


Moving in the direction w increases $g(x)$ fastest.



Changing w_0 moves the boundary parallel to itself.

$$\text{shift along } w = -\frac{\Delta w_0}{\|w\|}$$


A constant color offset can often be absorbed by the bias.



$g(x) = w^\top x + w_0$ changes if every parameter is rescaled.

$$d(x) = \frac{w^T x + w_0}{\|w\|}$$

$g(x)$ and $10g(x)$ define the same boundary but not the same raw magnitude.



A calculation at the sorting gate

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Teaching parameters

$$w = (1, -0.8)^T, w_0 = -3.5, x = (7.2, 4.1)^T$$

$$g(x) = 7.2 - 0.8 \times 4.1 - 3.5 = 0.42$$



Decision

The score is positive, so the model selects pack now.

Distance. $0.42 / \sqrt{1^2 + 0.8^2} \approx 0.328$ feature units.

Feature units change the weights

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Two equivalent records

Firmness 4.1 on a 0 to 10 scale equals 41 on a 0 to 100 scale.

Equivalent model

The weight -0.8 becomes -0.08 when that feature is multiplied by 10.

Wrong shortcut

Keeping the old weight after changing units moves the boundary.

Standardization makes scales comparable

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Transform each feature

$$z_j = \frac{x_j - \mu_j}{s_j}$$

Interpretation

A one unit change in z_j means one training standard deviation.

Deployment rule

Reuse the training values of μ_j and s_j for new samples.

The augmented vector absorbs the bias

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Augmented feature

$$\tilde{x} = (x_1, \dots, x_D, 1)^\top$$

Augmented weight

$$\tilde{w} = (w_1, \dots, w_D, w_0)^\top$$

$$g(x) = \tilde{w}^\top \tilde{x}$$

Why it helps

One inner product handles both feature weights and the bias.

Equivalent parameters and one boundary

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Positive rescaling


(w, w_0) and (cw, cw_0) define the same decision for every $c > 0$.

Negative rescaling

For $c < 0$, the boundary stays but the class labels swap.

Identifiability

A normalization or training objective selects one representative from many equivalent parameter sets.



There exists (w, w_0) such that $y_i(w^\top x_i + w_0) > 0$ for every sample.

One hyperplane places every positive sample on one side and every negative sample on the other.

Separable training points do not guarantee perfect predictions on future data.

x_1	x_2	Positive class
0	0	Yes
0	1	No
1	0	No
1	1	Yes

No affine function can place both diagonal corners on one side and the other corners on the opposite side.

Add an interaction feature, combine local lines, or use a deeper model.

One versus rest classification

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER


One score for each class

$g_c(x) > 0$ means class c versus all other classes.

$$g_c(x) = w_c^T x + w_{c0}, \quad c = 1, \dots, C$$

Three grades


Fit separate tests for premium, standard, and hold.



Several classes claim the same sample.

No classifier claims the sample.

Use comparable scores, a priority rule, or a reject option. State the choice explicitly.



Independent binary rules need a policy for overlaps and gaps.

One versus one classification

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER


Pairwise rule

Fit one classifier for every pair of classes.

$$\text{number of classifiers} = \frac{C(C-1)}{2}$$

Prediction

Each pair casts one vote. The class with most votes wins.



$C(C-1)/2$ pairwise classifiers; ties need a declared rule.

Pairwise voting can tie

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Three pairwise results

A beats B, B beats C, and C beats A.

Class	Votes
A	1
B	1
C	1



A complete classifier needs a tie rule

Examples include summed margins, calibrated pairwise probabilities, or a fixed class order.

One discriminant score per class

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER


Shared comparison

Fit C scores and choose the largest.

$$\hat{y} = \arg \max_c g_c(x), \quad g_c(x) = w_c^\top x + w_{c0}$$

Pairwise boundary

Classes c and k meet where $g_c(x) = g_k(x)$.



Argmax partitions the whole plane

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER


Complete assignment

Every point has at least one largest score.

Ties

Two or more classes can share the maximum exactly on a boundary.

Geometry. The visible class regions are intersections of linear inequalities.



Every point receives the class with the largest score; ties lie on boundaries.

A numeric three class decision

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Scores

$$g_1 = -x_1 + x_2, g_2 = x_1 + x_2 - 5, g_3 = -x_2 + 1$$

Input	g_1	g_2	g_3	Argmax
(6, 5)	-1	6	-4	Class 2
(3, 5)	2	3	-4	Class 2



One versus rest difference

At (3, 5), both g_1 and g_2 are positive.

Classifier count grows differently

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Classes	One versus rest	One versus one	Comparable class scores
3	3	3	3
10	10	45	10
100	100	4,950	100



Cost beyond training

Prediction must also evaluate the required scores or pairwise classifiers.

Scores and probabilities answer different questions

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Output	What it supports	What it does not guarantee
Score $g(x)$	Ranking and boundary decisions	Calibrated probability
Margin distance	Geometric separation	Frequency among similar cases
Probability $p(y x)$	Expected risk and thresholds	Correctness without validation

A link function can convert a score

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Logistic example

$$p(y = 1 | x) = \frac{1}{1 + \exp[-g(x)]}$$

Boundary

$g(x) = 0$ corresponds to probability 0.5.

Preview only

The link defines a probability model. Its fitting and evaluation appear in the linear models lecture.

Least squares can fit a discriminant

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Targets

Encode classes as $t_i \in \{-1, +1\}$.

Objective

$$\min_{\tilde{w}} \|\tilde{X}\tilde{w} - t\|_2^2$$

Closed form when appropriate

Solve $\tilde{X}^\top \tilde{X}\tilde{w} = \tilde{X}^\top t$.

Least squares uses every residual

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Correct but extreme

A positive sample with target $+1$ and score $+8$ has residual 7 .

Wrong but near

A positive sample with score -0.1 has residual -1.1 .

Mismatch


The first point is correctly classified but contributes a much larger squared error.

The fitted line follows the main cloud.

Squared loss gives the distant residual disproportionate influence.

Inspect leverage, measurement errors, and a classification specific loss.

Before the extreme point



Squared residual pulls the boundary



NumPy batch scoring

数字工程研究中心

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

```
import numpy as np

X = np.array([[7.2, 4.1],
              [5.4, 6.0],
              [6.8, 3.5]])
w = np.array([1.0, -0.8])
b = -3.5

scores = X @ w + b
labels = np.where(scores > 0, 1, -1)
print(scores)
print(labels)
```

Expected scores

[0.42, -2.90, 0.50] and labels [1, -1, 1].

NumPy multiclass scores

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

```
W = np.array([[ -1.,  1.],
               [ 1.,  1.],
               [ 0., -1.]])
b = np.array([0., -5., 1.])
X = np.array([[6., 5.], [3., 5.]])

scores = X @ W.T + b
prediction = scores.argmax(axis=1) + 1
print(scores)
print(prediction)
```

Axis check

Rows are samples and columns are class scores. Take the maximum across columns.

Transform x into features $\phi(x)$, then use a linear score in the transformed space.

$$g(x) = w^\top \phi(x)$$



A hyperplane in feature space can appear curved in the original space.

$$g(x) = w_0 + w_1\phi_1(x) + \cdots + w_M\phi_M(x)$$

If $\phi_1(x) = x^2$, the score is nonlinear in x .

Generalized linear here means a fixed feature map followed by a linear discriminant. It is distinct from the statistical family called generalized linear models.

One threshold cannot isolate a middle interval

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

One dimensional classes


Class 2 occupies $b < x < a$. Class 1 occupies $x \leq b$ or $x \geq a$.

Linear limitation

A one dimensional affine score changes sign at most once.

Needed boundary

The desired rule needs two crossing points.



Two zero crossings isolate the middle interval.



put into two features

$$\phi(x) = (x^2, x, 1)^\top$$

$$g(x) = (x - a)(x - b) = x^2 - (a + b)x + ab$$

$$w = (1, -(a+b), ab)^T$$

Decision. Positive outside the interval and negative inside when $a > b$.



The quadratic map turns two crossings on the line into one linear separator in feature space.

A line in feature space becomes a parabola

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Transformed coordinates

Let $z_1 = x^2$ and $z_2 = x$.

$$z_1 - (a + b)z_2 + ab = 0$$



Two views

This is a line in the z plane and a quadratic equation on the original x axis.

Exclusive or becomes linear with one interaction

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

x_1	x_2	$x_1 x_2$	Target y
-1	-1	+1	+1
-1	+1	-1	-1
+1	-1	-1	-1
+1	+1	+1	+1



A one feature rule

$$g(x) = x_1 x_2 \text{ and } \hat{y} = \text{sign } g(x).$$

What changed

The interaction records whether the two signs agree.

$$\phi(x) = (x_1^2, x_1x_2, x_2^2, x_1, x_2, 1)^\top$$

$$g(x) = w_{11}x_1^2 + w_{12}x_1x_2 + w_{22}x_2^2 + w_1x_1 + w_2x_2 + w_0$$

Ellipses, parabolas, hyperbolas, pairs of lines, and degenerate cases.

Quadratic feature count

Component	Number
Quadratic terms	$D(D + 1)/2$
Linear terms	D
Bias	1
Total	$(D + 1)(D + 2)/2$

For two inputs

$3 + 2 + 1 = 6$ features.

For one hundred inputs

5,151 features.

Polynomial degree grows the feature space

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

All monomials up to degree r

$$M = \binom{D+r}{r}$$

Input dimension D	Degree r	Features M
10	2	66
10	3	286
100	2	5,151
100	3	176,851



Capacity warning

Higher degree can overwhelm both memory and the available sample size.

Color ranges near 10 while acidity ranges near 0.1.

The scales become about 100 and 0.01.

Optimization and regularization can become dominated by feature units.

Usual response. Standardize suitable inputs before constructing polynomial terms.

x_1x_2 lets the effect of color depend on firmness.

Deep color may support pack now only when the tomato is not unusually hard.

A large coefficient can reflect chance correlation when data are scarce.

A richer map can also fit noise

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Training fit

A high degree polynomial can pass close to many training points.

Between the points

The boundary may oscillate in regions with little evidence.

Model selection

Use validation performance, regularization, and a degree chosen before looking at the test set.

Basis	Useful when	Main caution
Polynomial	Smooth global curvature	Feature count grows quickly
Radial basis	Local similarity matters	Centers and widths must be chosen
Spline	Smooth pieces along an ordered variable	Knots control flexibility
Domain interaction	A known mechanism combines features	Interpretation depends on preprocessing

```
import numpy as np

X = np.array([[1., 2.],
              [3., 4.]])
x1, x2 = X[:, 0], X[:, 1]
Phi = np.column_stack([
    x1**2, x1*x2, x2**2, x1, x2, np.ones(len(X))
])
print(Phi)
print(Phi.shape)
```

Two samples become two rows with six quadratic features.


```
scores = Phi @ w
labels = np.where(scores >= 0, 1, 2)
print(scores)
print(labels)
```

Class 1, class 2, class 2, class 1.

Use several affine pieces, each responsible for a local part of the input space.

Straight segments join to approximate a curved or multi region decision surface.

The number and placement of pieces must be learned or chosen.



Local affine pieces approximate a curved score while keeping each piece simple.

Several local lines fit a multimodal class

Two tomato varieties


The pack now class forms two separated clusters because the varieties mature at different colors.

One global line

It may cut through the hold cluster or miss one pack cluster.

Local pieces

Each cluster can receive a nearby linear comparison.



Two local separators handle a class with two separated modes.

Three boundary families on one problem

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Boundary family	Strength	Main limitation
One affine boundary	Simple and stable	Cannot bend
Feature mapped boundary	Smooth global curvature	Feature count may grow
Piecewise affine boundary	Local flexibility	Pieces and joins need control

No universal winner

Choose the smallest family that captures repeatable structure.

$$g_{c\ell}(x) = w_{c\ell}^\top x + b_{c\ell} \text{ for piece } \ell \text{ of class } c.$$
$$g_c(x) = \max_{\ell=1, \dots, L_c} g_{c\ell}(x)$$
$$\hat{y} = \arg \max_c g_c(x)$$

pieces meet on a line

Switching set

$$g_{c1}(x) = g_{c2}(x)$$

Meaning

The active piece changes across this hyperplane.

Not always a class boundary

Both pieces may belong to the same class.

Nearest prototypes create affine comparisons

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Choose the nearest prototype

$$\hat{j} = \arg \min_j \|x - \mu_j\|^2$$

$$-\|x - \mu_j\|^2 = -x^\top x + 2\mu_j^\top x - \|\mu_j\|^2$$

Shared term

$-x^\top x$ cancels when prototypes are compared.



Several prototypes bend the class boundary


Subclasses

Class c has prototypes $\mu_{c1}, \dots, \mu_{cL_c}$.

$$d_c(x) = \min_{\ell} \|x - \mu_{c\ell}\|^2, \quad \hat{y} = \arg \min_c d_c(x)$$

Geometry


Prototype Voronoi cells meet along line segments. Class labels merge neighboring cells.



Find one difficult local comparison between opposite classes.

Add another separating plane where errors remain.

End when validation improvement no longer justifies another piece.



Pieces	Likely behavior
1	A single affine boundary
A few	Captures broad local structure
Many	Tracks fine variations and noise
One per sample	Approaches nearest neighbor behavior

Choose the piece count with validation evidence, not training error alone.

pieces underfit 工程研究中心

Systematic errors occur in the same region for many samples.

The local class structure is missing from the model.

Useful response. Add a justified prototype or feature, then recheck on validation data.


ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER



A local area where samples or prototypes from different classes are close.

Use nearby opposite class evidence to design the local separator.

Easy interior points still matter for defining where the local rule is allowed to act.




ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

A prototype from each class is the other's nearest opposite class prototype.

$$\left[x - \frac{\mu_a + \mu_b}{2} \right]^\top (\mu_a - \mu_b) = 0$$

Use the perpendicular bisector, then refine with nearby samples.




A perpendicular bisector gives an initial local plane for each opposite-class prototype pair.

The distance matrix has three samples by four prototypes.

Project many features onto one axis while preserving class separation.

$$Z = W^T X$$


Projected class centers should be far apart and each class should remain compact.



Only the direction of projection changes.

Different classes land on overlapping positions.

Class centers separate and within class spread stays small.




Projection changes the one-dimensional overlap, not the original points.

The projected histograms overlap strongly.

The axis preserved a direction with large within class variation.


Lesson. Separation in the original plane does not survive every one dimensional view.



The projected histograms occupy different parts of the axis.

The direction keeps center separation while reducing within class spread.

Real classes may still overlap after the best linear projection.




Projection onto $z = w^T x$

If $\|w\| = 1$, then $z = w^\top x = \|x\| \cos \theta$.

Replacing w by cw rescales every projected value by c .

Geometry. The perpendicular foot on the axis determines the scalar coordinate.



Original mean

$$m_k = \frac{1}{N_k} \sum_{i: y_i=k} x_i$$

Projected mean

$$\tilde{m}_k = \frac{1}{N_k} \sum_{i: y_i=k} w^\top x_i = w^\top m_k$$

Consequence

Projected center difference is $w^\top (m_1 - m_2)$.


Class scatter

$$S_k = \sum_{i: y_i=k} (x_i - m_k)(x_i - m_k)^\top$$

Projected scatter

$$\tilde{S}_k^2 = \sum_{i: y_i = k} (w^\top x_i - w^\top m_k)^2 = w^\top S_k w$$

Total within scatter

 $w^\top S_W w$ where $S_W = S_1 + S_2$.

Between class separation

Scalar separation

$$(\tilde{m}_1 - \tilde{m}_2)^2 = [w^\top (m_1 - m_2)]^2$$

Matrix form

$$S_B = (m_1 - m_2)(m_1 - m_2)^\top$$

$$[w^\top (m_1 - m_2)]^2 = w^\top S_B w$$

Projected class means should be far apart.

Projected samples within each class should stay compact.

$$J(w) = \frac{w^\top S_B w}{w^\top S_W w} = \frac{[w^\top (m_1 - m_2)]^2}{w^\top S_W w}$$

The Fisher ratio ignores scale

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Rescale the direction

Let $w' = cw$ for any nonzero c .

$$J(cw) = \frac{c^2 w^\top S_B w}{c^2 w^\top S_W w} = J(w)$$




Practical normalization

Use $w/\|w\|$ for display, or fix $w^\top S_W w = 1$ in a derivation.

$$W = m_1 - m_2$$

It ignores directions where each class has large variance.

The center line may align with a noisy long axis and create overlapping projections.



Maximize $w^\top S_B w$ subject to $w^\top S_W w = 1$.

$$S_B W = \lambda S_W W$$

Two class simplification. $S_B w = (m_1 - m_2)(m_1 - m_2)^\top w$ is a scalar multiple of $m_1 - m_2$.

$$w \propto S_W^{-1}(m_1 - m_2)$$

$$w \propto S_W^{-1}(m_1 - m_2)$$

The inverse scatter reduces its contribution.

A modest mean difference can matter when measurements are consistent.

Centering all samples by the global mean mixes between class separation into the within class scatter.

$$w = S_W^{-1}(m_1 - m_2)$$

Solve $S_W w = m_1 - m_2$.

```
Sw = np.array([[4., 1.], [1., 2.]])
mean1 = np.array([1., 1.])
mean2 = np.array([3., 2.])
w = np.linalg.solve(Sw, mean1 - mean2)
```

Explicit matrix inversion is slower and usually less numerically stable.

A two dimensional Fisher example

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Class means

$$m_1 = (2, 0)^T \text{ and } m_2 = (2, 2)^T$$

Within class scatters

$$S_1 = \begin{bmatrix} 1 & 1/2 \\ 1/2 & 1 \end{bmatrix} \text{ and } S_2 = \begin{bmatrix} 1 & -1/2 \\ -1/2 & 1 \end{bmatrix}$$

$$S_W = 2I, \quad w = S_W^{-1}(m_1 - m_2) = (0, -1)^T$$

Projected scores in the numeric example

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Point	$z = w^T x$	Side of midpoint
(2, 0)	0	Class 1 center
(2, 2)	-2	Class 2 center
(4, 0.4)	-0.4	Class 1 side
(1, 1.6)	-1.6	Class 2 side



Midpoint threshold

$t = (0 - 2)/2 = -1.$

A threshold is chosen after the direction

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Projection

Fisher selects w to separate one dimensional class distributions.

Classification

Choose a threshold t and compare $z = w^\top x$ with t .

Equal symmetric case

The midpoint $(w^\top m_1 + w^\top m_2)/2$ is a common choice.

Separate decisions

The Fisher criterion alone does not encode priors or unequal error costs.

Class priors move the projected threshold

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Equal projected variance

If $z | k \sim \mathcal{N}(\mu_k, \sigma^2)$, Bayes gives

$$t = \frac{\mu_1 + \mu_2}{2} - \frac{\sigma^2}{\mu_1 - \mu_2} \log \frac{\pi_1}{\pi_2}$$



Interpretation

A more common class receives a larger region along the projection axis.



$$x \mid y = k \sim \mathcal{N}(\mu_k, \Sigma)$$

$$W_{\text{Gaussian}} = \Sigma^{-1}(\mu_1 - \mu_2)$$

Fisher uses $S_W^{-1}(m_1 - m_2)$, which has the same form.

Priors and losses still determine the intercept of the Gaussian decision rule.

Features may be duplicated, constant, or more numerous than independent centered samples.

Thousands of pixels with only dozens of labeled images cannot produce a full rank within class scatter.

A direct solve fails or becomes extremely sensitive to small perturbations.

Regularized scatter

$$S_W^{(\lambda)} = S_W + \lambda I, \quad \lambda > 0$$

Direction

$$w_\lambda = (S_W + \lambda I)^{-1}(m_1 - m_2)$$

Effect

Small variance directions no longer receive arbitrarily large inverse weights.

Selection

Choose lambda with validation data.

Model checks 数字工程研究中心

- 1 Compare training and validation separation on the projected axis.
- 2 Inspect projected distributions rather than only their means.
- 3 Report the threshold, priors, and loss assumptions separately.
- 4 Stress test the direction after small data and scaling changes.

A high criterion is not enough

A few outliers or a shifted deployment population can still produce poor decisions.

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

Design	What is learned	Best first question
Linear	one score $w^\top x + w_0$	Can one straight boundary work?
Feature map	a linear score in $\phi(x)$	Which transformation exposes the pattern?
Piecewise linear	several local linear scores	Does the boundary need several faces?
Fisher	a projection with low within class spread	Which direction separates compactly?

Begin with the simplest geometry that can express the decision.

Size separates most samples, but two growing regions create a bent boundary in color and firmness.

- 1 Plot the two features and try one linear score as a baseline.
- 2 Add interpretable transformed features if the bend has a known form.
- 3 Use local linear pieces if different growing regions form separate subgroups.
- 4 Use Fisher projection when a compact one dimensional inspection score is required.



Choose the threshold from the cost of sending a wrong variety to packaging.

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

We designed decision functions and understood their geometry.

The perceptron will learn a separating direction from mistakes; decision trees will partition the space by questions.

A learning rule is not only an optimizer. It also determines which boundaries are easy to discover.

Exit ticket 先进计算与数字工程研究中心

ADVANCED COMPUTING AND DIGITAL TECHNOLOGY RESEARCH CENTER

- 1 State the difference between a generative and a discriminative classifier in one sentence.
- 2 Explain the sign and magnitude of a linear discriminant score.
- 3 Give one reason to use a feature map and one reason to use linear pieces.
- 4 Explain the numerator and denominator of the Fisher criterion.
- 5 For XOR, choose a model family and justify the choice.



A strong answer

Names the geometry, the information used, and one failure mode rather than listing formulas.